

Final orange count per episode

LightwheelAI
Baseline

Teleop
Monotask

Teleop
Subtask

Teleop+Auto
Monotask


Teleop+Auto
Subtask

Standard fine-tuning

only the action expert and state projection trained; language model and vision encoder frozen

SmoVLA: batch 64, 20k steps unless noted

Mean /3




Partial fine-tuning

language model also trained; vision encoder still frozen

batch 64, 20k steps unless noted

Mean /3

Share of episodes (%)




Full fine-tuning

all weights trained, vision encoder included

all models: batch 32, 40k steps

Mean /3



3/3 oranges
2/3 oranges
1/3 oranges
0/3 oranges